

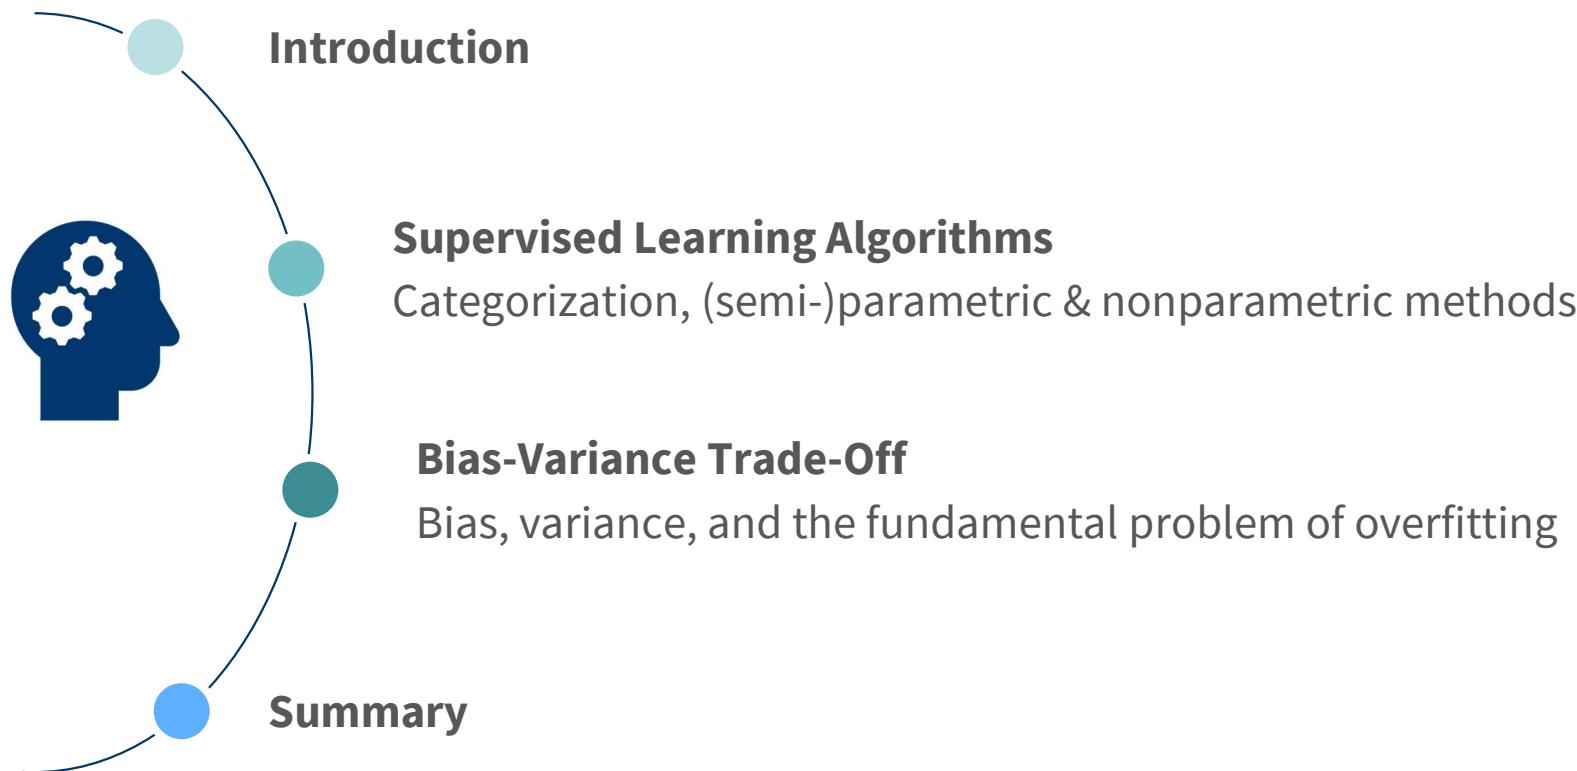


VHB ProDok – Machine Learning – Block I

L.2: Foundations of Supervised Learning

Stefan Lessmann

Agenda

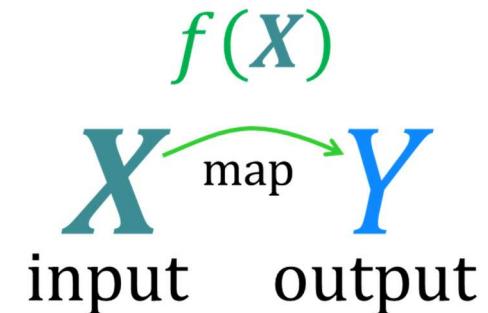




Introduction

Supervised Machine Learning (SML)

Input to output mapping



- **The output variable defines a prediction target**

- For an individual subject, this is the value we wish to know
 - We cannot observe the target (at the right time) or face a high cost to measure it (e.g., human labor)

- **The (input) features characterize each subject**

- These values are easy/cheap to observe
 - We believe they determine the target value

- **SML algorithms approximate the hypothesized, unknown relationship between feature values and the target**

- They receive data that exemplifies the relationship in the form of pairs of feature and target values
 - Using their internal logic, they then learn the mapping functions from these examples

Formalization of Supervised Machine Learning

Resale price forecasting example

- We aim at forecasting resale prices (our target variable) denoted by y
- We assume that resale prices y depend on features x
 - We do not know how exactly resale prices depend on feature values
 - But we have access to historical data $\mathcal{D} = \{y_i, x_i\}_{i=1}^n$ that exemplifies the relationship
- At decision time, i.e., when need a forecast, we can observe x but not y
- We use algorithms to learn a model f that maps from features to target $f: x \mapsto y$

PRODUCT	LIST PRICE [\$]	AGE [month]	CLIENT INDUSTRY	...
Dell XPS 15'	2,500	36	Mining	...
Dell XPS 15'	2,500	24	Health	...
Dell XPS 17'	3,000	36	Manufacturing	...
Lenovo Yoga 11'	799	12	Office	...
Lenovo Yoga 13'	1,100	12	Office	...
...

$$x = (x_1, x_2, \dots, x_m) \in \mathbb{R}^m$$

$$f(x)$$

OBSERVED RESALE PRICE [\$]
347
416
538
...
266
...

$$y \in \mathbb{R}$$

Two Flavors of Supervised Machine Learning

Regression & classification methods model numerical & discrete targets

■ Credit scoring example

- Financial institutions receive loan applications $\mathbf{x} \in \mathbb{R}^m$
- Loan approval decisions depend on an estimate of the applicant's probability to default $y \in \{0; 1\}$
- When deciding on an application, we **can observe x but not y**
- Historical data $\mathcal{D} = \{\mathbf{y}_i, \mathbf{x}_i\}_{i=1}^n$ illustrates which granted applications were repaid / defaulted

■ We use algorithms to learn a model f that maps from features to target $f: \mathbf{x} \mapsto y$

BUREAU SCORE	DEBT/ INCOME	COLLATERAL	YEARS AT ADDRESS	AGE	...
650	20%	Yes	2	<21	...
280	43%	No	0	21-29	...
$\mathbf{x} = (x_1, x_2, \dots, x_m) \in \mathbb{R}^m$					
575	33%	No	12	>50	...
715	24%	Yes	1	21-29	...
580	18%	No	6	40-50	...

$$f(\mathbf{x})$$

REPAYMENT STATUS
Good (0)
Bad (1)

$$y \in \{0; 1\}$$

Good (0)
Good (0)
Bad(1)

Credit Scoring: Predictive Analytics Perspective

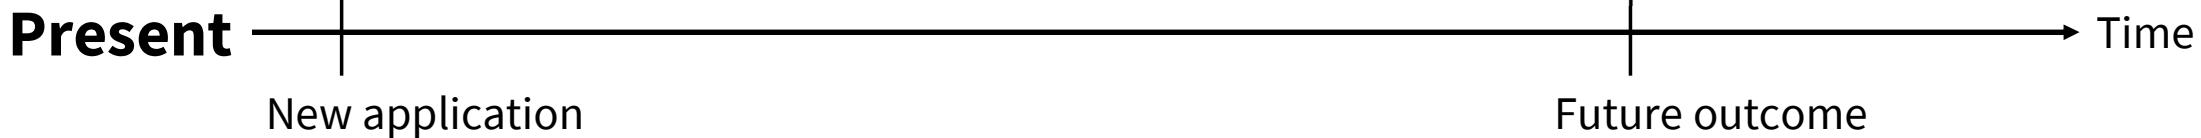
Train model on historical data to predict the risk of a new application



CLIENT	BUREAU SCORE	DEBT/INCOME	COLLATERAL	LOAN/VALUE	YEARS AT ADDRESS	AGE	...
Mary	650	20%	Yes	45%	2	<21	...
Peter	280	43%	No	55%	0	21-29	...
Ying	750	27%	Yes	89%	8	30-39	...
Carlos	600	18%	Yes	70%	4	40-50	...
Emre	575	33%	No	82%	12	>50	...
Güley	715	24%	Yes	62%	1	21-29	...
Shih-Chi	580	18%	No	78%	6	40-50	...
Alice	800	14%	Yes	37%	10	40-50	...
Sash	410	29%	Yes	105%	4	21-29	...

$$f: x \mapsto y$$

REPAYMENT STATUS
GOOD
BAD
GOOD
BAD
GOOD
BAD



BUREAU SCORE	DEBT/INCOME	...	REPAYMENT STATUS
735	25%	...	???

$$f(x) = \hat{y}$$

PREDICTION
LIKELY GOOD

Two-Stage SML Paradigm

Generalize from training data to unseen data

Stage 1: Model Training



Data-driven development of a predictive model using **labelled data** $\mathcal{D} = \{\mathbf{y}_i, \mathbf{x}_i\}_{i=1}^n$

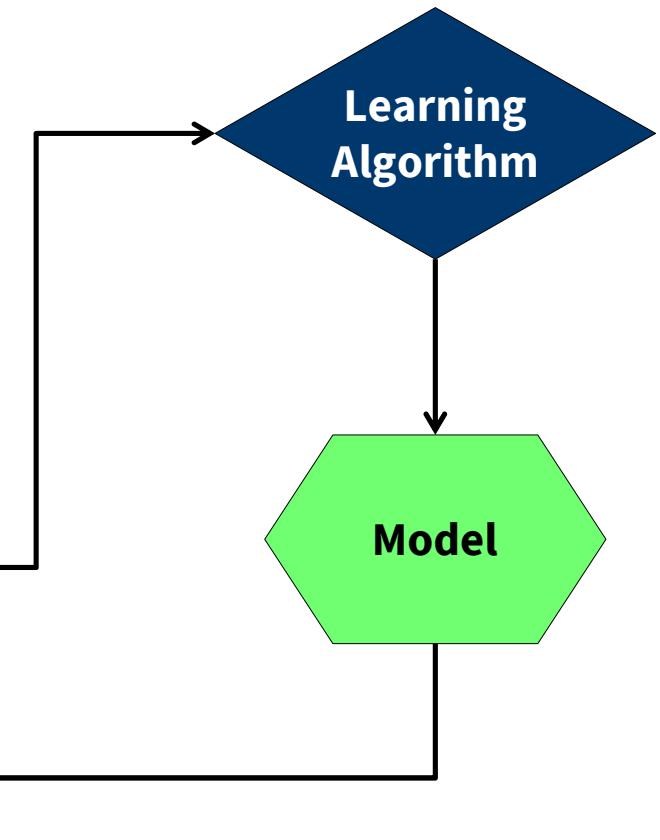
Training data incl. Y					
i	X_1	X_2	...	X_m	Y
1
2
...
n

Stage 2: Model Application

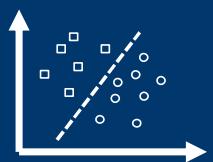


Application of **trained model** to unseen data (only **features**) to obtain a **prediction**

New data w/o Y				
i	X_1	X_2	...	X_m
$n + 1$
$n + 2$
...
N



i	\hat{y}
$n + 1$...
$n + 2$...
...	...
N	...



Supervised Learning Algorithms

Categorization, (semi-)parametric & nonparametric
methods

Algorithms for Supervised Learning (Selection)

Approaches split into (semi-)parametric and nonparametric models

■ Parametric models assume a functional form for the feature-target relationship

- The model is fully characterized by a fixed number of parameters θ
 - Leading to the stochastic model equation $y = f(x; \theta) + \epsilon$
 - With ϵ denoting the error term and $\theta \in \mathbb{R}^k, k < \infty$

■ Nonparametric models learn the functional from data

■ Semi-parametric models: parametric structure + nonparametric components

Tree- and prototype-based algorithms
(non-parametric)

- CART, CHAID, C4.5
- Bagging / Random Forest
- Gradient Boosting / XGB
- Nearest neighbors
- ...

Regression-type algorithms
(semi-/parametric)

- Generalized linear models (GLM)
- Generalized additive models (GAM)
- Artificial neural networks (ANN)
- Support vector machines (SVM)
- ...

Linear Models for Supervised Machine Learning

Three well-known econometric approaches

- **Conditional mean structure**

$$\mathbf{y} = \mathbb{E}[Y|\mathbf{X} = \mathbf{x}] + \epsilon$$

- **Additive predictor**

$$\eta(\mathbf{x}) = b + \sum_{j=1}^m f_j(\mathbf{x}_j)$$

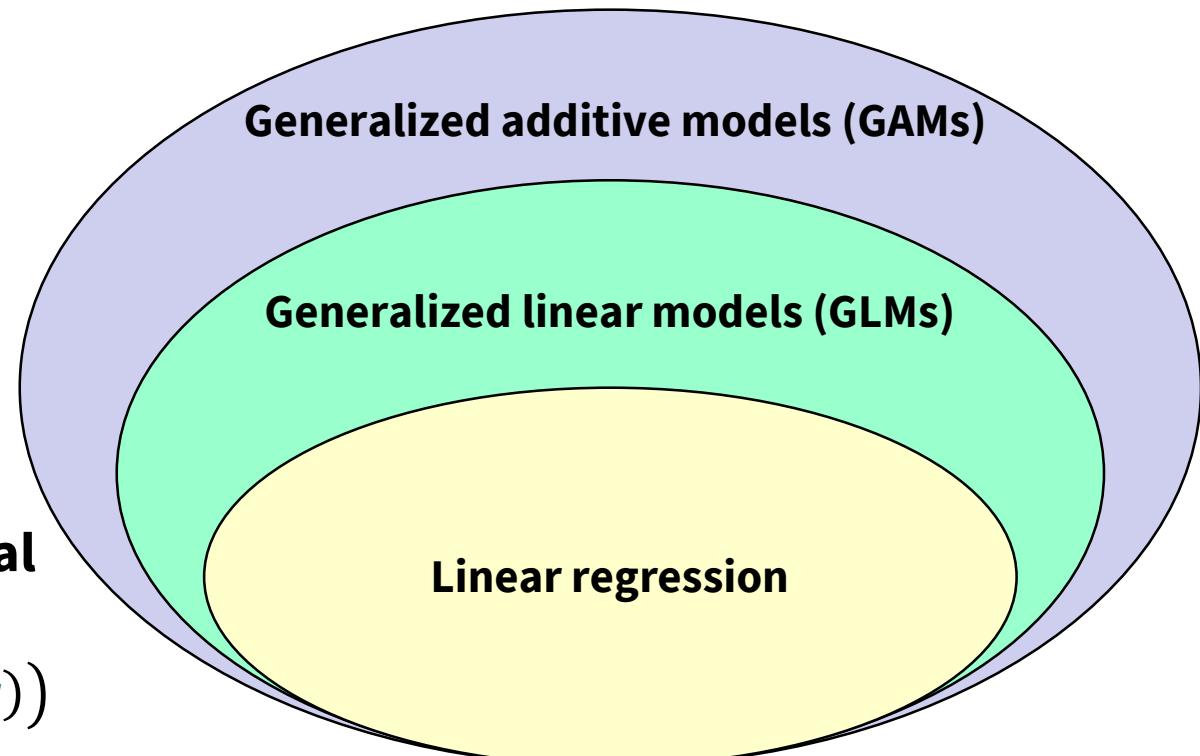
- With f_j denoting some smooth function
- For GLM $f_j(\mathbf{x}_j) = w_j \mathbf{x}_j$

- **Link function g connecting conditional expectation to additive predictor**

$$g(\mathbb{E}[y|\mathbf{x}]) = \eta(\mathbf{x}) \Leftrightarrow \mathbb{E}[y|\mathbf{x}] = g^{-1}(\eta(\mathbf{x}))$$

- **Prediction equation**

$$\hat{y} = g^{-1}(\eta(\mathbf{x}))$$



Linear Models for Supervised Machine Learning

Three well-known econometric approaches

■ Linear regression: linearity in conditional mean and Gaussian errors

- Model equation: $\mathbb{E}[Y|X = \mathbf{x}] = \eta(\mathbf{x}) = \mathbf{b} + \sum_{j=1}^m w_j x_j$
- and $g(\cdot) = \text{id}$

■ Generalized linear models (GLMs) add a link function for non-Gaussian targets

- $g(\mathbb{E}[Y|X = \mathbf{x}]) = \eta(\mathbf{x}) = \mathbf{b} + \sum_{j=1}^m w_j x_j$
- Example: logistic regression for binary outcomes, $g := \log\left(\frac{p}{1-p}\right)$ with $p = \Pr(Y = 1 | \mathbf{x})$

■ Generalized additive models (GAMs): replace linear effects by smooth nonlinear functions while keeping additivity and interpretability

- Model $g(\mathbb{E}[Y|X = \mathbf{x}]) = \mathbf{b} + \sum_{j=1}^m w_j f_j(x_j)$
- Link function depends on the target's distribution (exactly as in GLM)
- Binary → logit, count data → log, Gaussian → identity

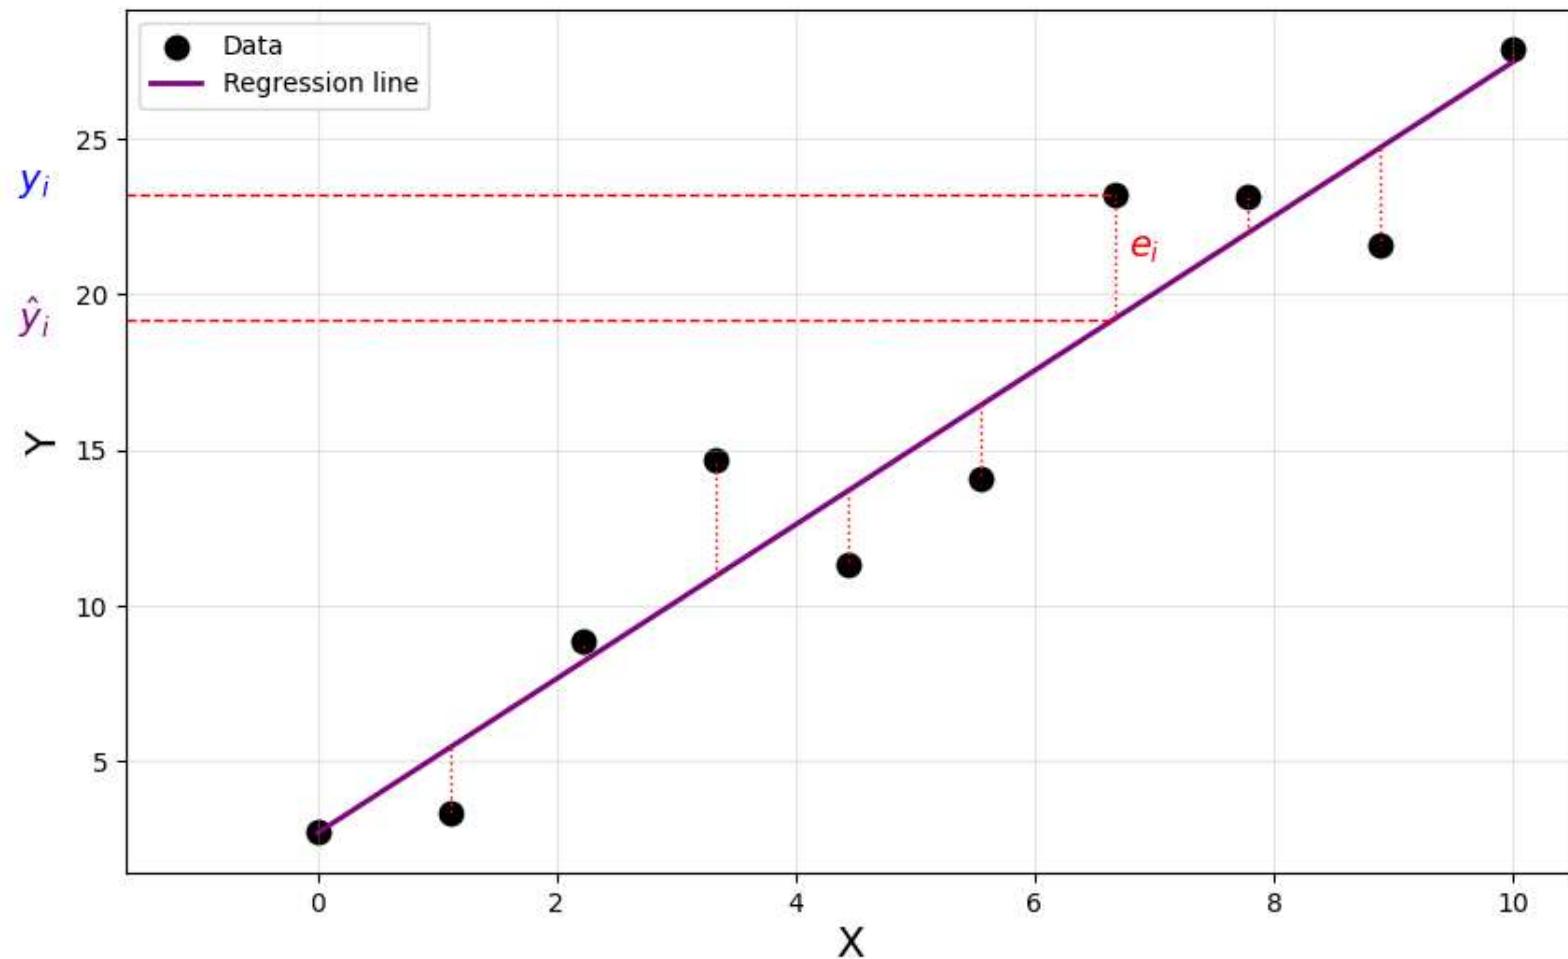
Empirical Risk Minimization (ERM) Principle

Train (aka estimate) model through minimizing a loss function

- Residuals e

- Loss function

- Loss minimization



Empirical Risk Minimization (ERM) Principle

Train (aka estimate) model through minimizing a loss function

- Residuals e measures deviation between actual target value and model prediction

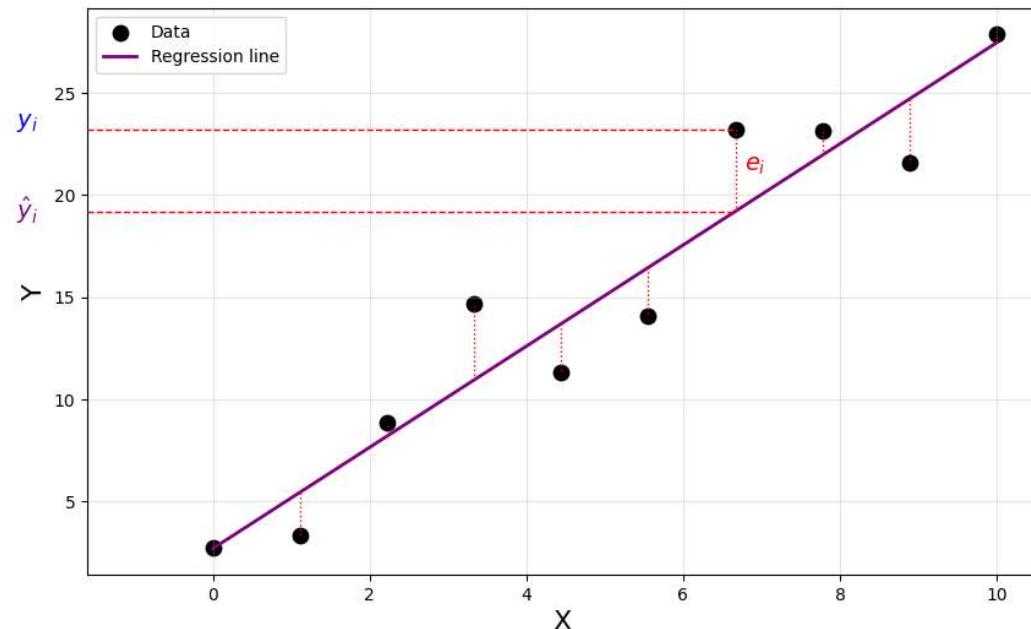
$$e = (y - \hat{y})$$

- Loss function assesses model fit

- Squared error loss function
 - Pointwise loss $J(y_i, \hat{y}_i) = e_i^2 = (y_i - \hat{y}_i)^2$
 - Dataset-level loss $J(y, \hat{y}) = \sum_{i=1}^n e_i^2$
- Several other loss functions exist

- Loss minimization

- OLS estimator $\hat{b}, \hat{w} = \operatorname{argmin}_{(\mathbf{b}, \mathbf{w})} \frac{1}{n} \sum_{i=1}^n (y_i - (\mathbf{b} + \sum_{j=1}^m w_j x_j))^2$
- OLS estimator has an analytical solution (i.e., *normal equation*)
- In general, we perform the minimization using iterative algorithms (e.g., gradient descent)



Empirical Risk Minimization (ERM) Principle

Estimating generalizable models by minimizing empirical risk

■ Minimize the expected loss of a model over the joint distribution of \mathbf{X} and \mathbf{Y}

- Population loss $J(\mathbf{f}) = \mathbb{E}_{(X,Y)} [J(\mathbf{Y}, \mathbf{f}(\mathbf{X}))]$
- ERM objective $\mathbf{f}^* = \operatorname{argmin}_{\mathbf{f} \in \mathcal{F}} \mathbb{E}_{(X,Y)} [J(\mathbf{Y}, \mathbf{f}(\mathbf{X}))]$

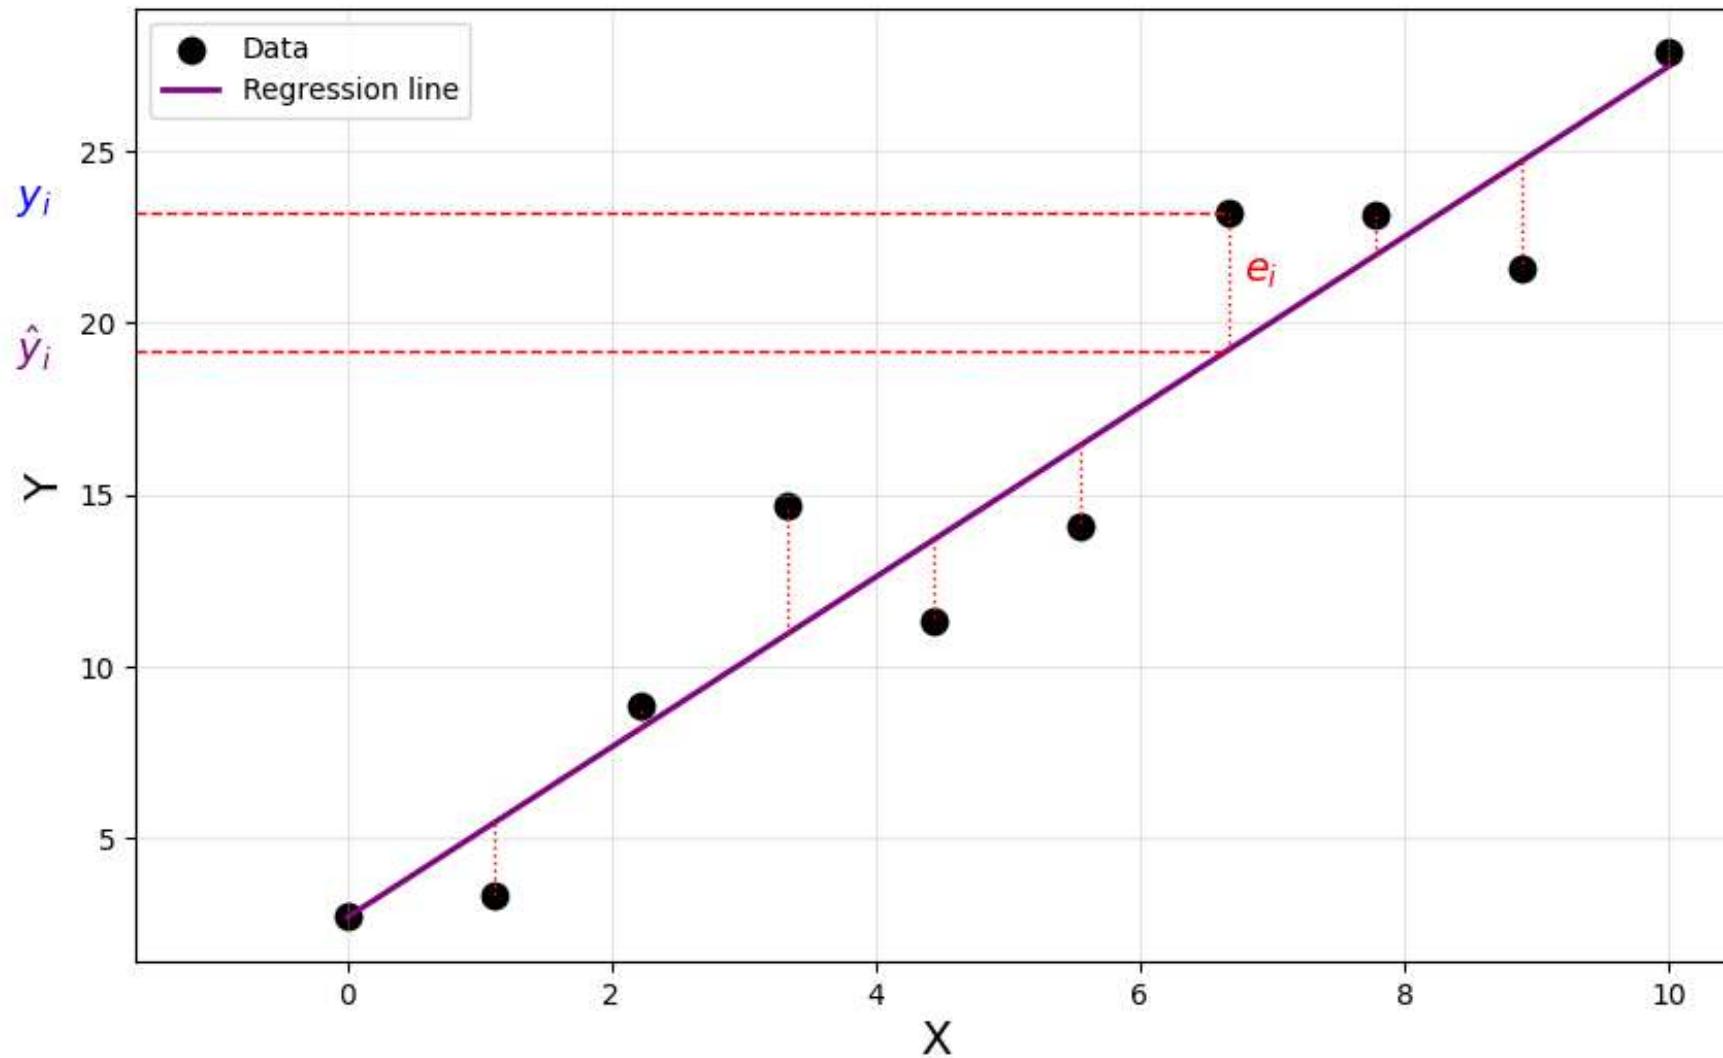
■ Empirical risk

- Loss computed over a dataset $\mathcal{D} = \{\mathbf{y}_i, \mathbf{x}_i\}_{i=1}^n$
- Empirical loss/risk $\hat{J}(\mathbf{f}) = \frac{1}{n} \sum_{i=1}^n J(\mathbf{y}_i, \mathbf{f}(\mathbf{x}_i))$
- ERM estimator $\hat{\mathbf{f}} = \operatorname{argmin}_{\mathbf{f} \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n J(\mathbf{y}_i, \mathbf{f}(\mathbf{x}_i))$
- with \mathcal{F} denoting the class of functions the model can learn

■ Maximum likelihood estimation

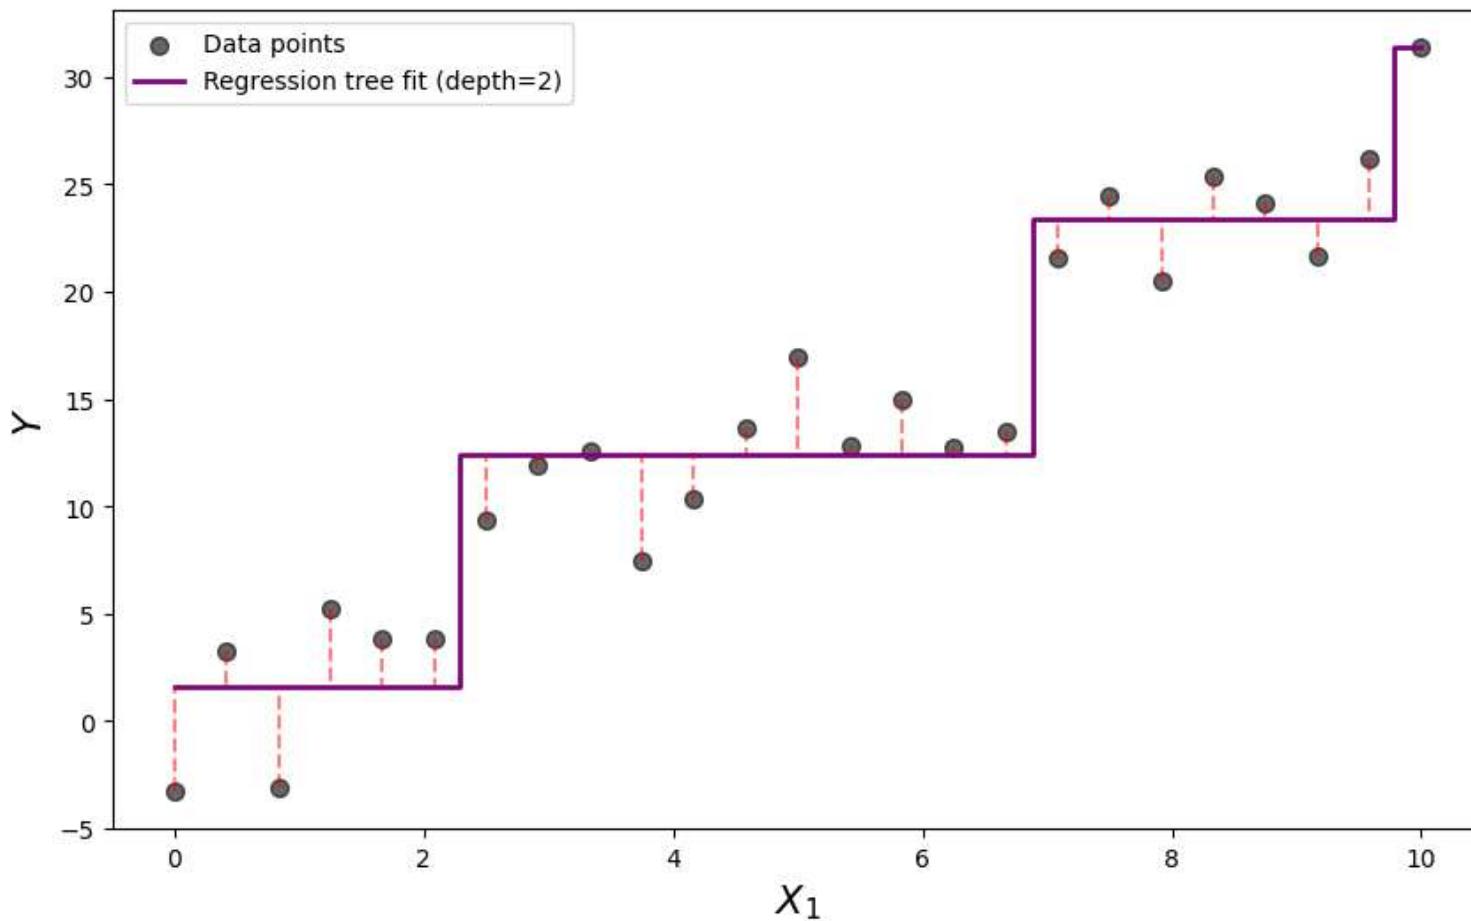
- Idea is to maximize the probability of the model given the data
- Specific instance of the ERM principle

Excusus: Loss Minimization and Probability Maximization



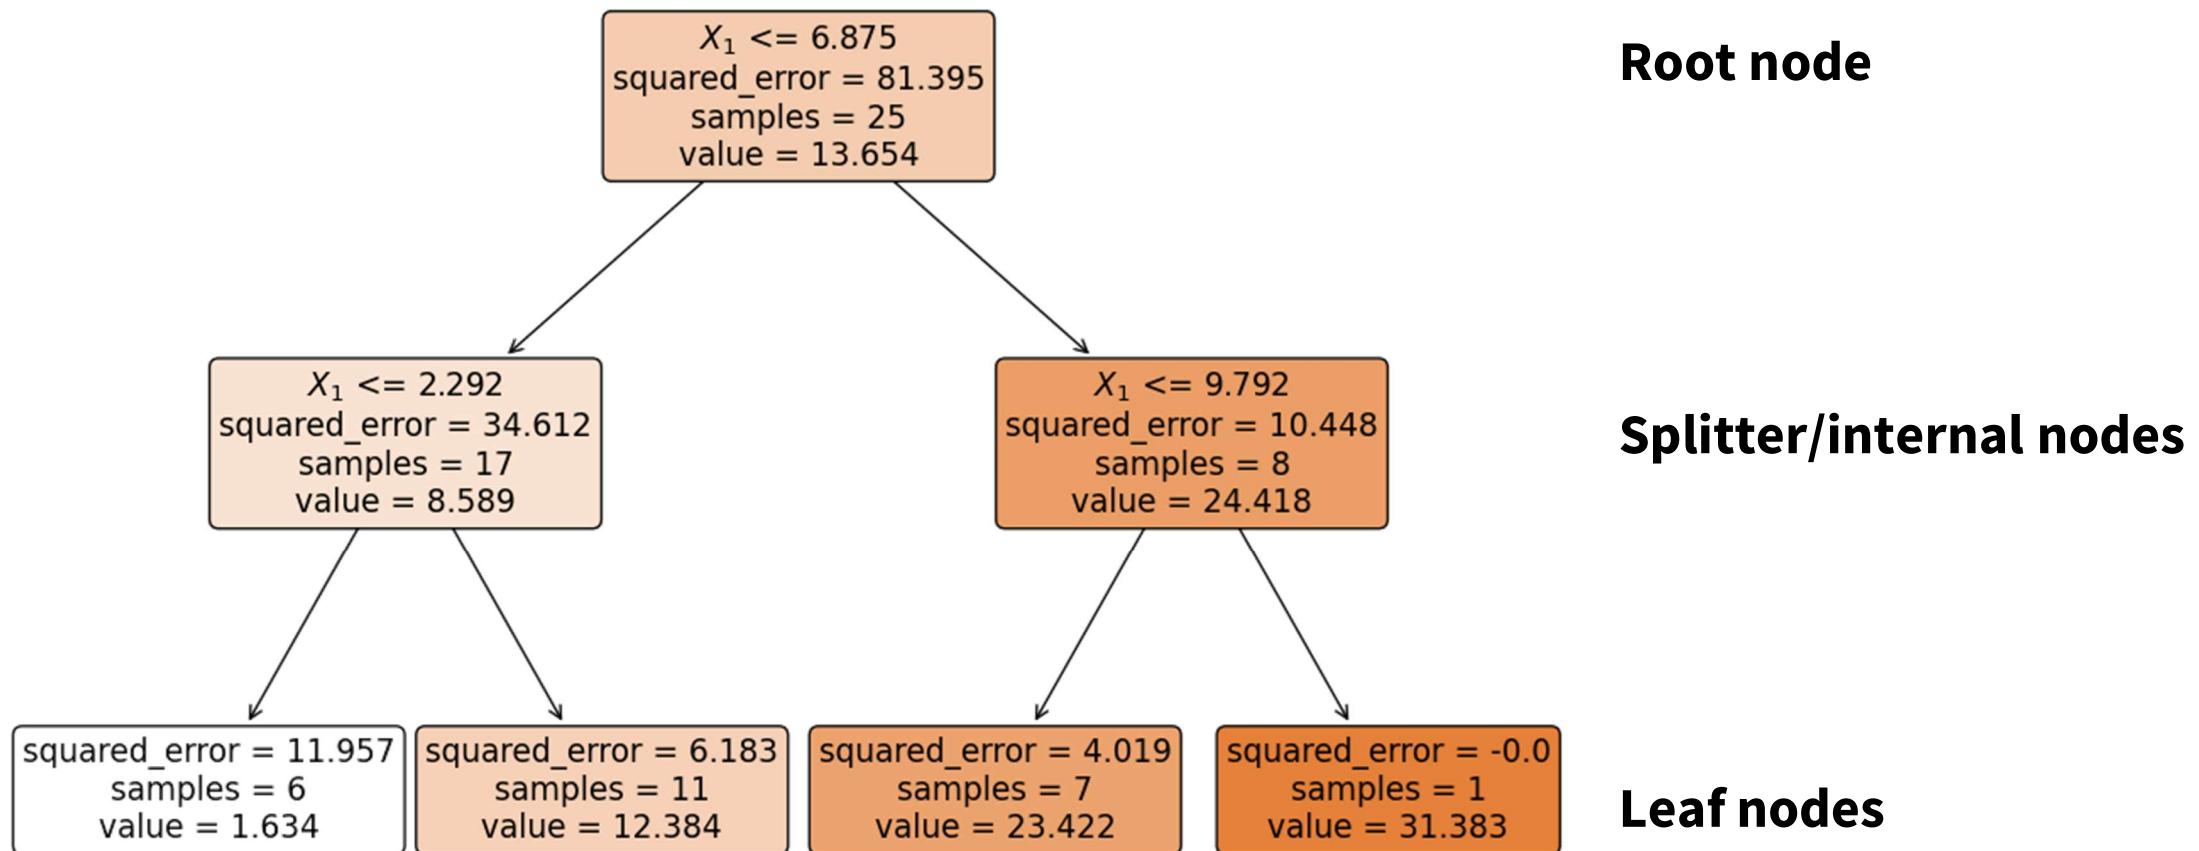
Beyond (Semi) Parametric Models: Classification & Regression Trees

Tree-based methods recursively partition the feature space



Beyond (Semi) Parametric Models: Classification & Regression Trees

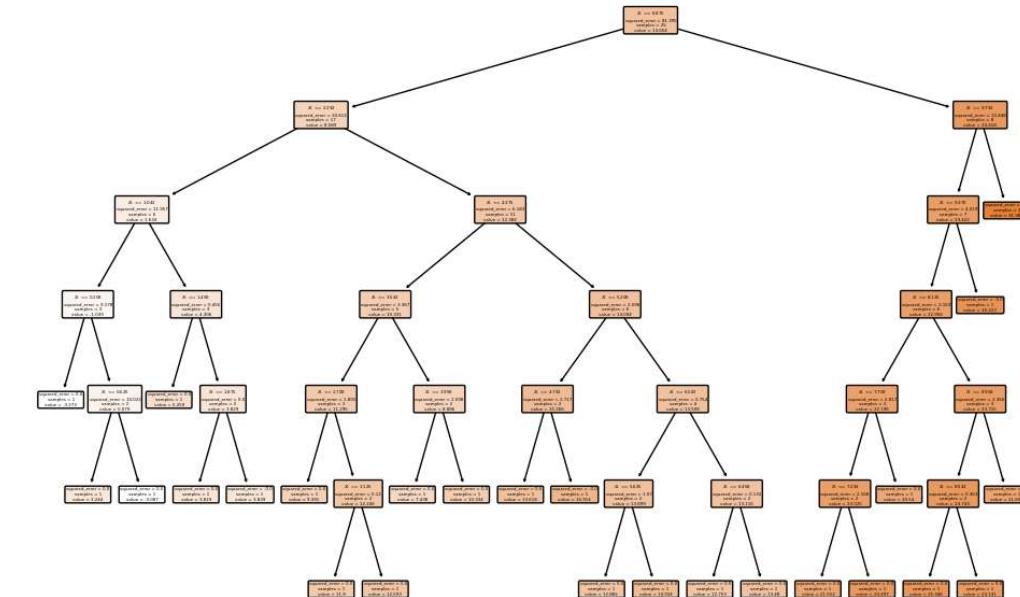
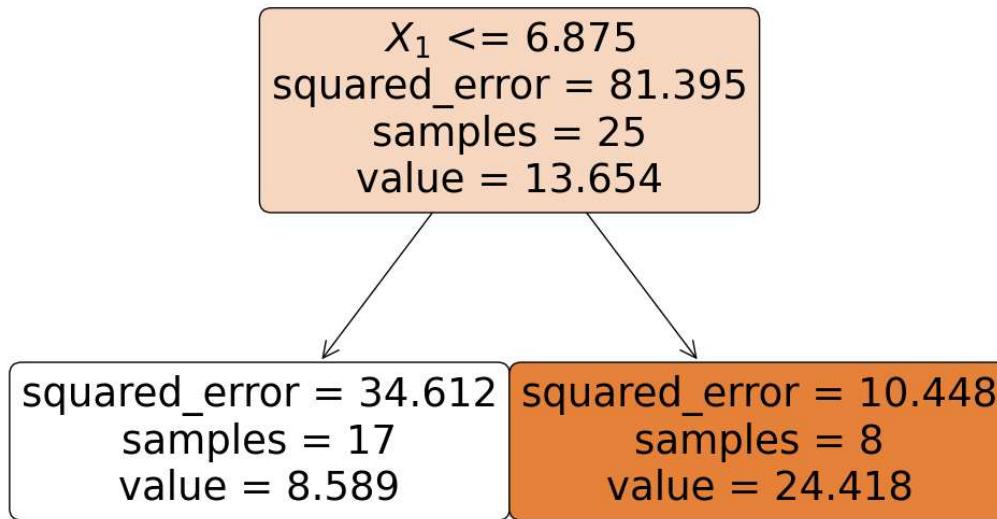
Tree-based methods recursively partition the feature space



Simple and Complex Trees

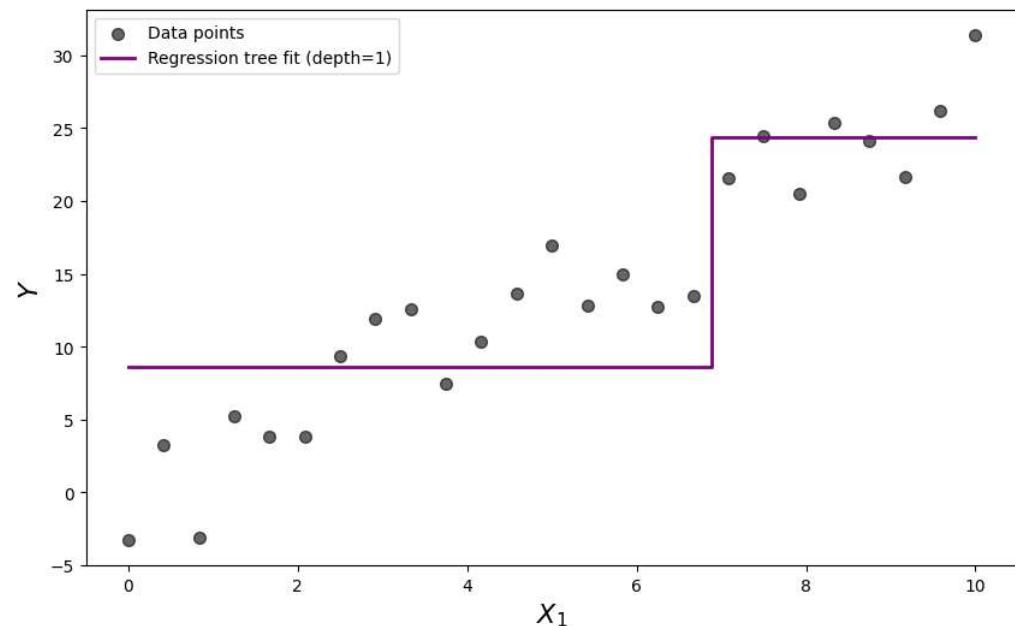
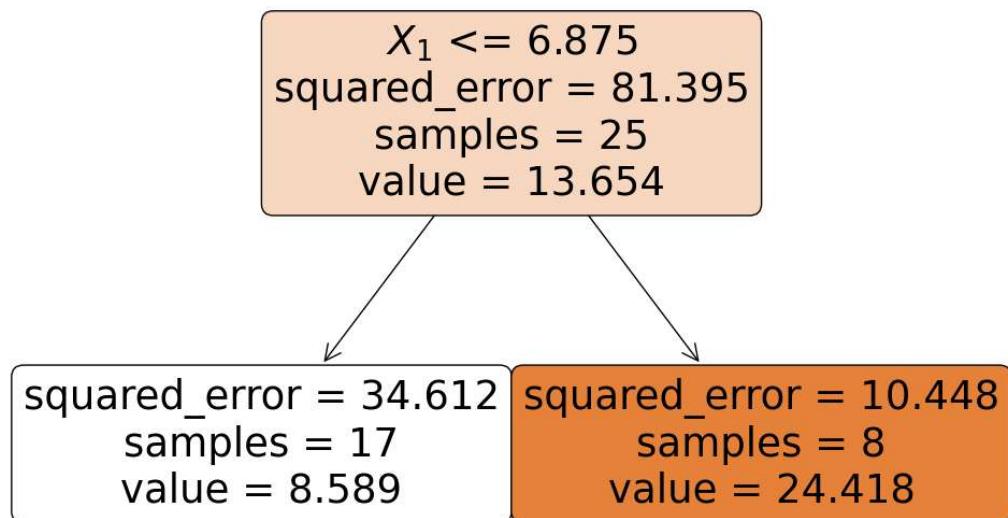
Tree depth governs model complexity. Complexity determines residuals.

- Hyperparameters govern the behavior of ML algorithms (e.g., depth of a tree)
- Many hyperparameters concern the complexity of the learnt model
- Model complexity is closely related to model residuals



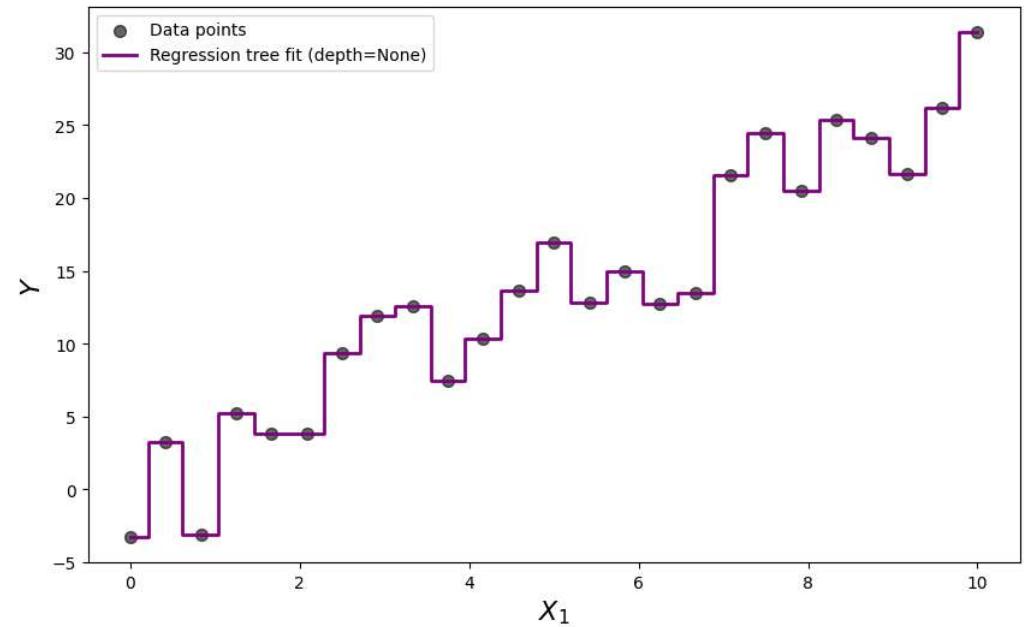
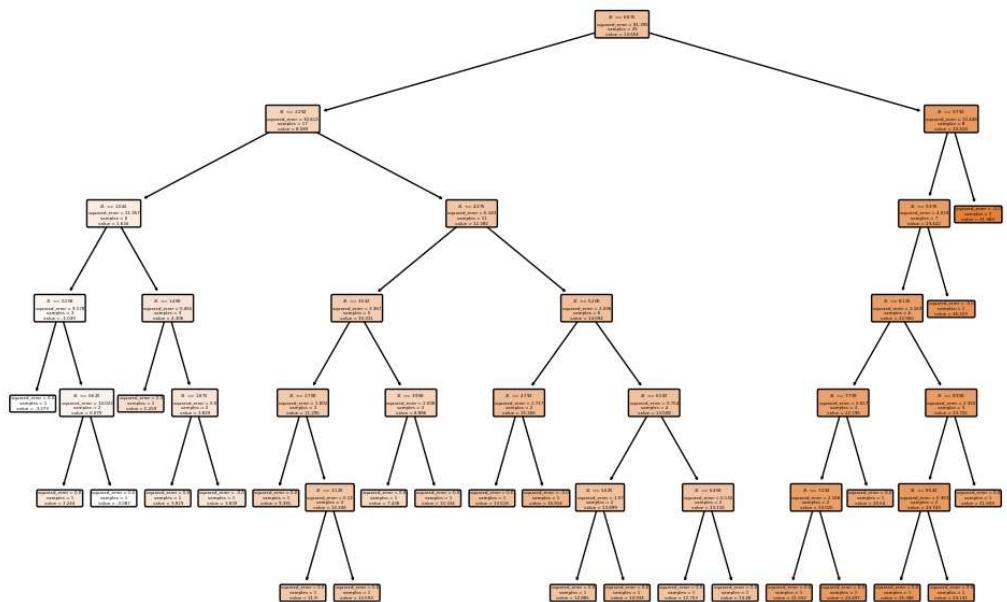
Simple and Complex Trees

Total sum of (squared) residuals is large for the shallow tree



Simple and Complex Trees

Residuals are zero for the deep tree



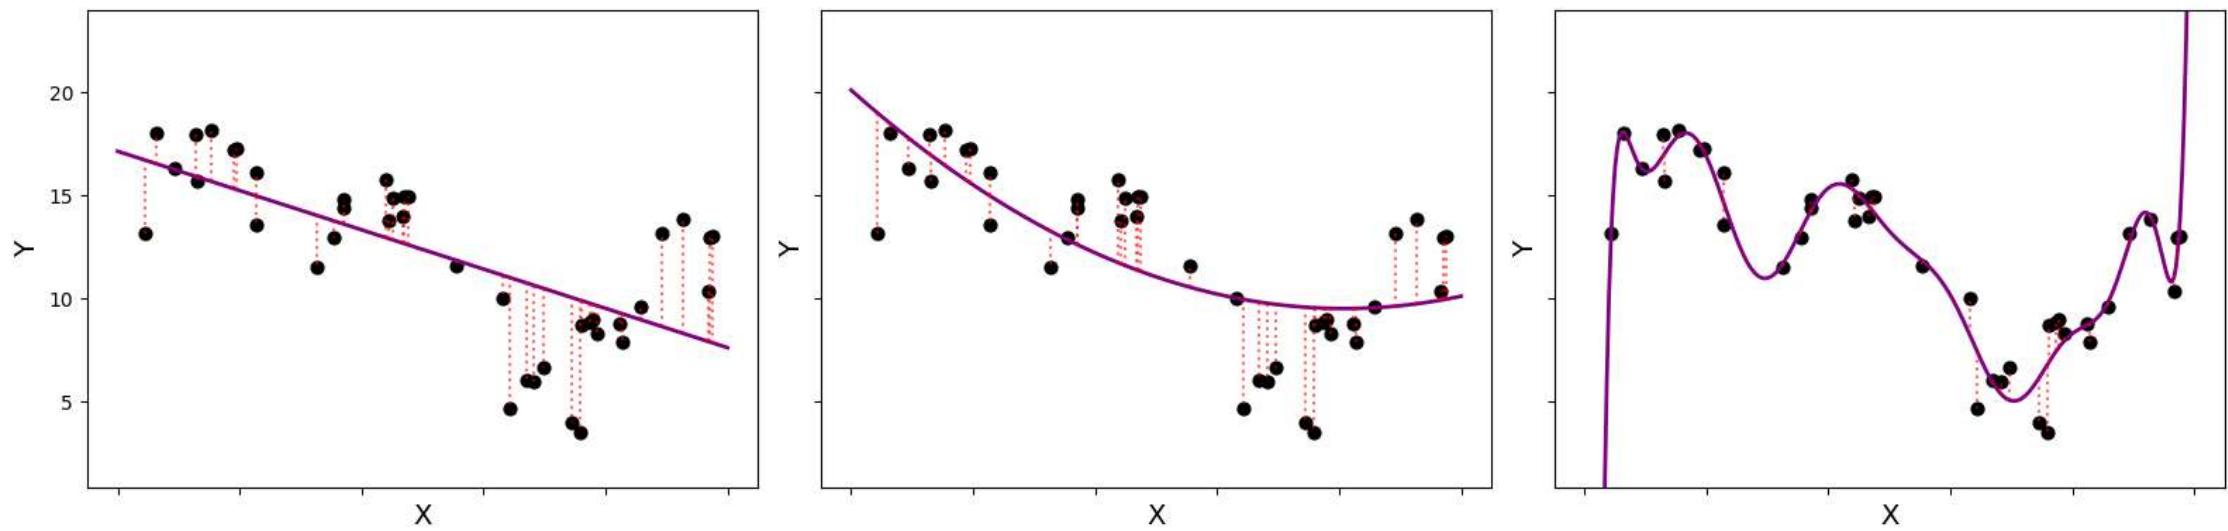


Bias-Variance Trade-Off

Bias, variance, and the fundamental problem of overfitting

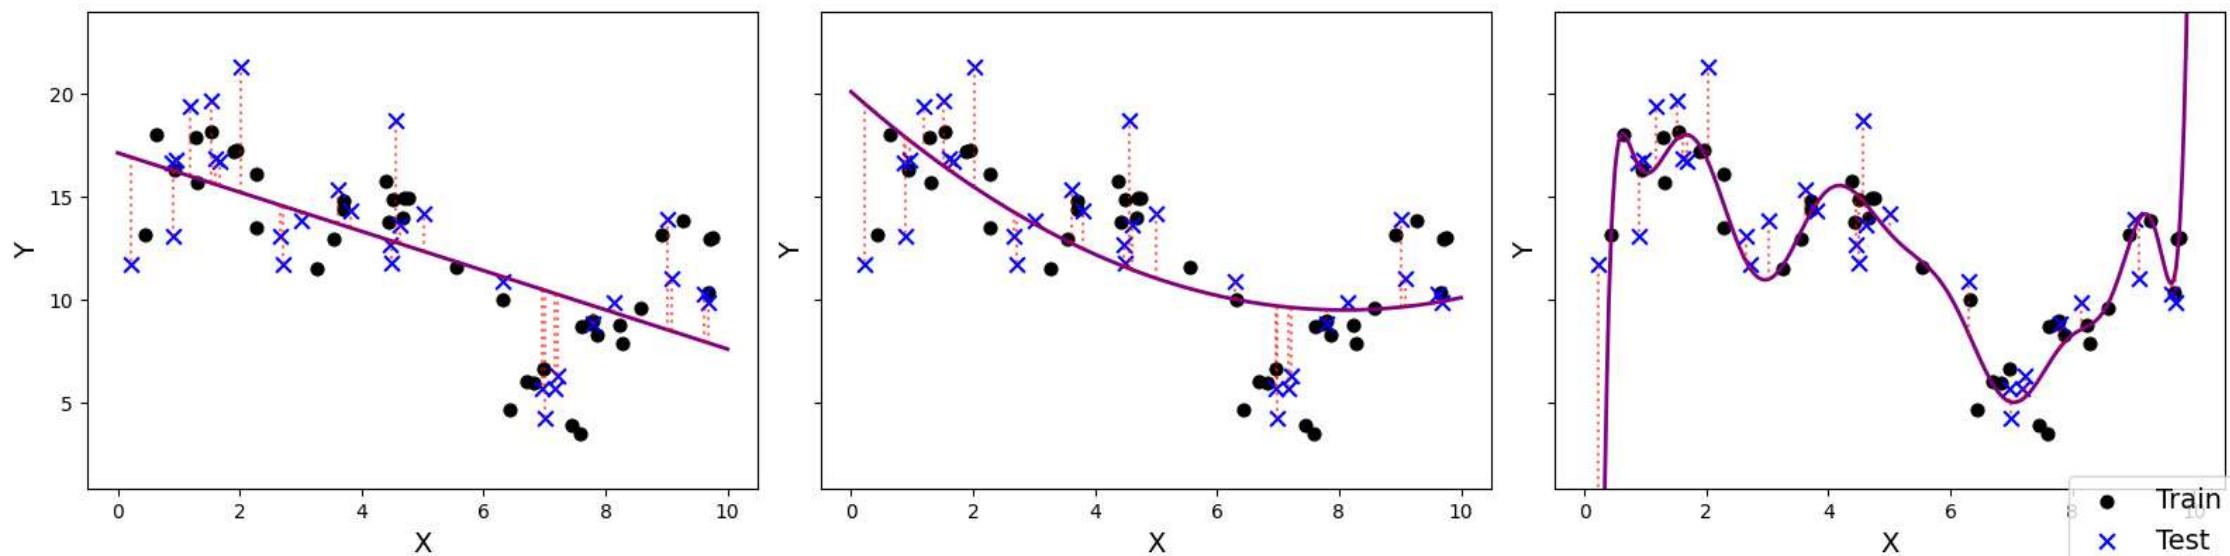
Three Candidate Regression Models for a Given Dataset

Which model do you prefer?



Three Candidate Regression Models for a Given Dataset

ML models can overfit the training data leading to poor generalization



The Fundamental Problem of Overfitting

But why would a model overfit?

■ Recall how we organize model training

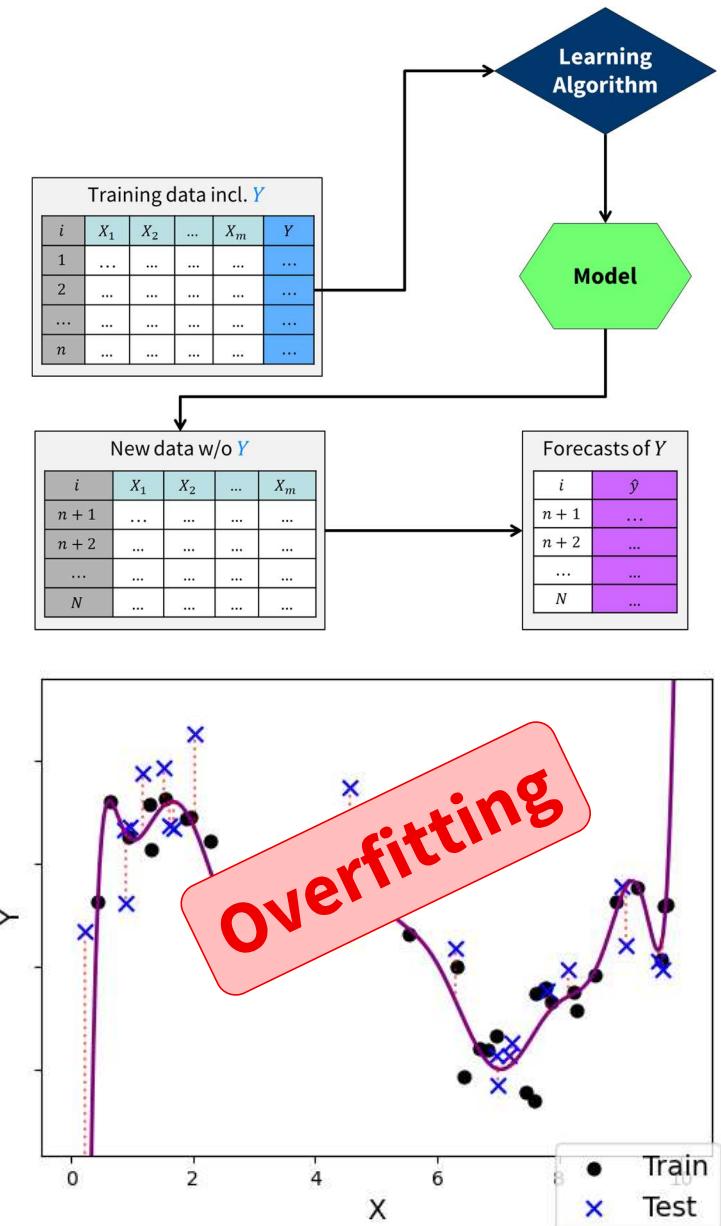
- Loss minimization → minimal loss is *mathematically* optimal
- *Optimal* fit of the training sample does not imply a good model (i.e., capable to generalize to new data)
- The model might be too specific and geared toward the training sample

■ The training data set is a sample

- We assume the sample represents the population well
- But every sample exhibit random (sampling) variation

■ A complex ML model can capture both, true structure and idiosyncrasies of training sample

■ Large error if applying such model to new data



The Fundamental Problem of Overfitting

But why would a model overfit?

■ Recall how we organize model training

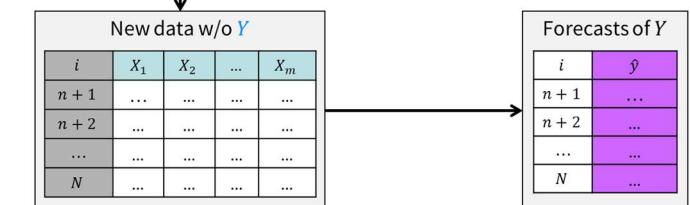
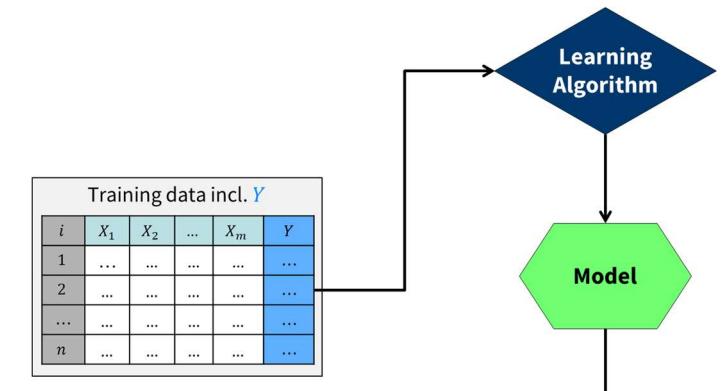
- Loss minimization → minimal loss is *mathematically* optimal
- *Optimal* fit of the training sample does not imply a good model (i.e., capable to generalize to new data)
- The model might be too specific and geared toward the training sample

■ The training data set is a sample

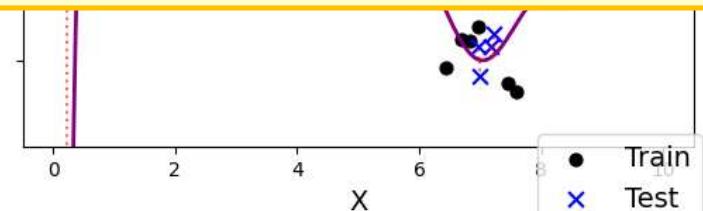
- We assume the sample represents the population well
- But every sample exhibit random (sampling) variation

■ A complex ML model can capture both, true structure and idiosyncrasies of training sample

■ Large error if applying such model to new data



But why is overfitting an ML problem? Do statistical models also overfit?



The Trade-Off Between Bias and Variance

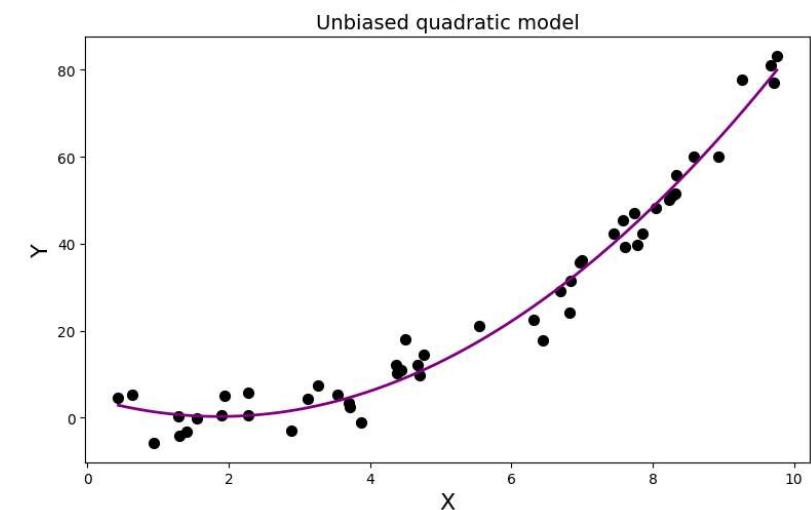
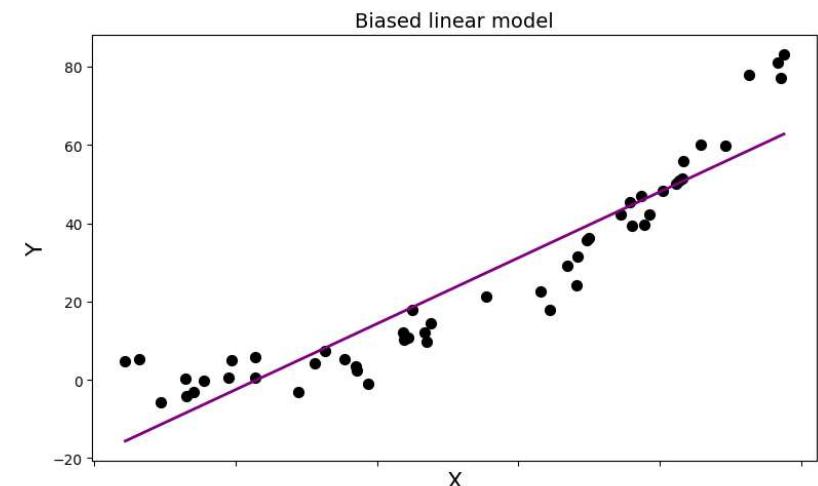
The generalization error depends on a model's bias and variance

■ Generalization error: a model's error on *any* data from the population

- True target of supervised machine learning
- Not observable in practice
- Approximated by the empirical loss on a data sample

■ Bias: can the model learn the relationship between features and the target?

- Recall ERM objective
$$f^* = \operatorname{argmin}_{f \in \mathcal{F}} \mathbb{E}_{(X,Y)}[J(Y, f(X))]$$
- Bias asks whether the algorithm can learn
$$f^*(x) = \mathbb{E}[Y|X = x]$$
- The more complex a model the lower its bias

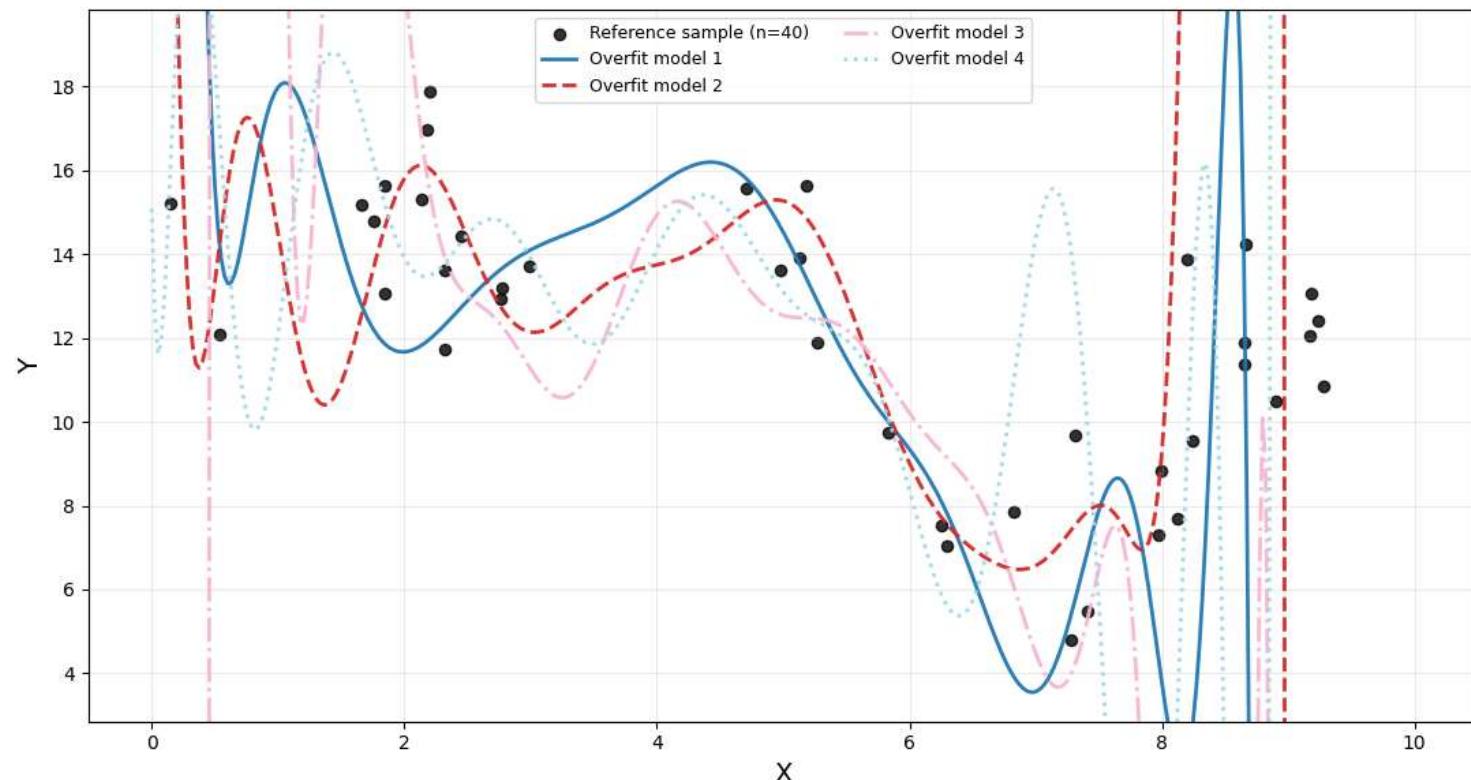


The Trade-Off Between Bias and Variance

The generalization error depends on a model's bias and variance

- Generalization error: a model's error on *any* data from the population
- Variance: how sensitive is the model to the training sample?

- Assume you estimate a linear regression model over many random samples drawn from one dataset
- How much will **coefficients** (i.e., the **model**) vary across samples?
- Changes in **models** imply changes in forecast,
- So variance \approx variation in **model predictions**



Bias-Variance Trade-Off and Overfitting

Overfitting is a problem of complex (low-bias) models that suffer high variance

- **Model complexity is about its expressive power**

- **Simple models**

- High bias & low variance
 - E.g., linear regression

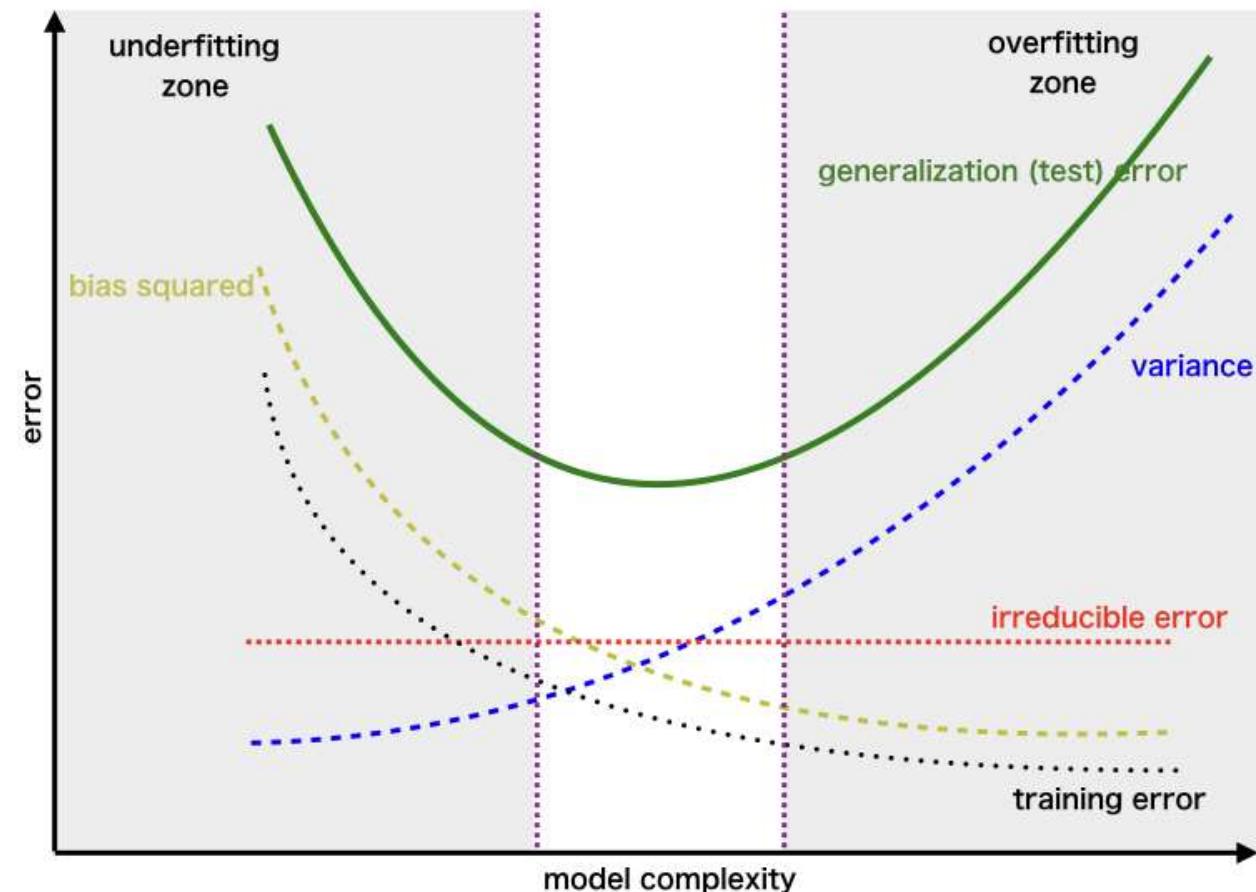
- **Complex models**

- Low bias & high variance
 - E.g., deep regression tree

- **SML aims at a good compromise**

- **Common paradigm in ML**

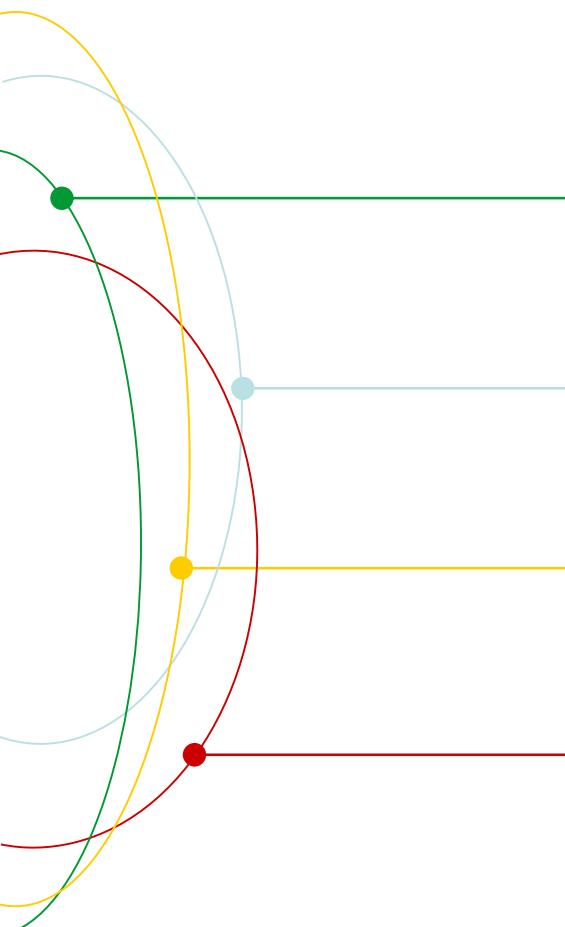
- Use an advanced, complex model
 - Manage / control complexity somehow





Summary

Summary



Learning goals

- Supervised machine learning paradigm
- Bias-variance trade-off and overfitting



Findings

- SML approximates a functional mapping from inputs to outputs
 - Parametric models assume a functional form a prior
 - Nonparametric models learn the mapping from data
- Training involves minimizing a loss function
- Complex models can overfit the training data leading to poor generalization
- Overfitting is a problem of low bias and high variance



What next

- Tree-based SML algorithm
- Practices for model evaluation

Thank you for your attention!

Stefan Lessmann

Chair of Information Systems
School of Business and Economics
Humboldt-University of Berlin, Germany

Tel. +49.30.2093.5742
Fax. +49.30.2093.5741

stefan.lessmann@hu-berlin.de
<http://bit.ly/hu-wi>

www.hu-berlin.de

